

Exercise 1

(a)

```
In [5]: # Imports
import numpy as np
import matplotlib.pyplot as plt
from sklearn.neighbors import NearestNeighbors
import time

In [6]: # 1) Load data
data_path = 'data/digit_features_normalized.npy'
labels_path = 'data/digit_labels.npy'
X = np.load(data_path)
labels = np.load(labels_path) if __import__('os').path.exists(labels_path) else None
print(X.shape)
if labels is not None:
    print('Labels shape try to align them with samples. If labels length matches columns, transpose X.')
    if labels.shape[0] == X.shape[1]:
        print('Labels match number of rows -> using rows as samples')
    else:
        print('Labels match number of columns -> transposing X to have samples as rows')
        X = X.T
    print('new X shape:', X.shape)
else:
    print('Warning: labels length does not match rows or columns of X; ignoring labels')
    labels = None
print('Number of samples N =', N)
raw X shape: (16, 2333)
raw labels shape: (2233)
Labels match number of columns -> transposing X to have samples as rows
raw X shape: (2233, 16)
raw labels shape: (2233, 1)
Number of samples N = 2233
Number of samples N = 2233

In [7]: # 2) Build symmetrized kNN graph (k=15) and extract edges E
k = 15
nbrs = NearestNeighbors(n_neighbors=k, algorithm='auto', metric='euclidean').fit(X)
dist, inds = nbrs.kneighbors(X)
N = X.shape[0]
E = np.zeros((N, N))
if not np.all(np.isfinite(dist)):
    print('Warning: adjacency then symmetrize')
    adj = np.zeros(N, N, dtype=bool)
    for i in range(N):
        for j in range(i+1, N):
            if dist[i][j] < k:
                adj[i][j] = True
    E = adj + adj.T
    E[E > 1] = 1
else:
    E[inds[:, :k], np.arange(N)] = 1
    E[inds[:, :k].T, np.arange(N)] = 1
    print('Number of nodes N =', N)
    print('Number of edges |E| =', E.shape[0])
    print('Number of nodes N = 2233')
    print('Number of edges |E| = 45955')

In [8]: # 3) Sample random repulsive pairs: |R| = 50
def sample_random_pairs(R, mask):
    """Return arrays (a,b) of shape (R,) with a != b
    If rng is None:
        -> R = np.random.default_rng()
    R = int(factor * N)
    a = rng.integers(0, N, size=R, endpoint=False)
    b = rng.integers(0, N, size=R, endpoint=False)
    # avoid self-pairs by resampling those entries
    mask = np.ones(N, dtype=bool)
    while mask.any():
        b[mask] = rng.integers(0, N, size=mask.sum(), endpoint=False)
    mask[a == b] = False
    return a, b

# quick test
a, b = sample_random_repulsive_pairs(100, factor=5)
print('sampled repulsive pairs:', a.shape[0])
sampled repulsive pairs: 500

In [9]: # 4) Parameters and initialization
# alpha = 10.0 # repulsion strength
seed = 42
rng = np.random.default_rng(seed)
# Initialize 2D embeddings
scale = 1.0
Y = rng.normal(scale=scale, size=(N, 2))
print('Initialized embeddings Y shape:', Y.shape)
Initialized embeddings Y shape: (2233, 2)

In [10]: # 6) Force (gradient) computation using analytical derivatives
def compute_grad(Y, E, rep_a, rep_b, c):
    """Compute gradient of loss function w.r.t. (i,j) pairs, rep_a, rep_b: (R, R)
    Y: (N,2), E: (N,2) = {(i,j)} pairs, rep_a,rep_b: (R, R)
    N = Y.shape[0]
    grad = np.zeros_like(Y)
    grad -= attractive term over edges E
    if E.shape[0] == 0:
        return grad
    i = E[:, 0]
    j = E[:, 1]
    diff = Y[i] - Y[j] # (R,2)
    d1 = np.sum(diff, axis=1) # (R,1)
    denom = 1.0 + d1 * d1 # (R,1)
    deno2 = 1.0 + d2 * d2 # (R,1)
    # gradient contribution: 2*(y_i-y_j)/(1+d1^2) for i; symmetric for j
    coef = (-2.0 * c / denom)[i, None] # (R,1)
    contrib = coef * diff # (R,2)
    grad += contrib
    # np.add.at(grad, i, contrib)
    np.add.at(grad, j, -contrib) # opposite sign for j
    # avoid self-pairs, a contrib
    np.add.at(grad, a, -contrib)
    np.add.at(grad, b, -contrib) # symmetric
    return grad

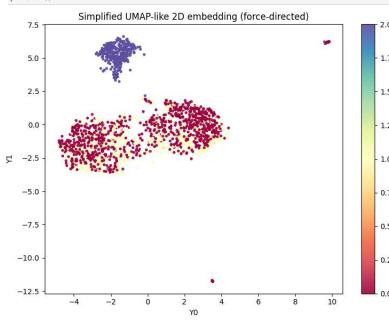
In [11]: # 7) Optimization: gradient descent with decreasing learning rate
def optimize(Y, E, N, c, rng, n_iter=300, lr=0.5, rep_factor=5, verbose=True):
    Y = compute(Y)
    history = {'loss': []}
    for it in range(n_iter):
        lr = lr * (1 - it / float(n_iter))
        Y = lr * Y + (1 - lr * float(it)) * Y
        # Compute learning-rate decay
        if it % 10 == 0:
            print(f'it {it} loss {history["loss"][-1]}')
        # Compute gradients
        a, b = sample_random_repulsive_pairs(N, rep_factor, rng=rng)
        grad = compute_gradients(Y, E, a, b, c)
        # Compute gradients (gradient descent)
        Y -= lr * grad
        # Compute loss (optional) every 10 iters
        if it % 10 == 0:
            if it < n_iter - 1:
                # attractive loss
                if E.shape[0] > 0:
                    a = rep_a
                    b = rep_b
                    d1r = np.sum(Y[a] - Y[b], axis=1) # (R,1)
                    d2r = np.sum(d1r * d1r, axis=1) # (R,1)
                    deno2r = 1.0 + d2r * d2r # (R,1)
                    coefr = (-2.0 * c / deno2r)[i, None] # (R,1)
                    contribr = coefr * d1r # (R,2)
                    contribr = contribr / deno2r # (R,2)
                    grad -= contribr
                loss = loss_attr + loss.rep
            history['loss'].append(it, loss_attr, loss.rep, loss)
        if verbose:
            print(f'it {it} (it:{it}) {lr} loss_attr:{loss_attr:e} loss_rep:{loss.rep:e} total:{loss:e}')
    return Y, history

# Run optimization (keep iterations modest for demo; increase as needed)
start = time.time()
Y_opt, history = optimize(Y, E, N, c, rng, n_iter=300, lr=0.5, rep_factor=5, verbose=True)
print('Optimization time (s):', time.time() - start)
it 0/300 lr=0.5000 loss_attr=2.732e+04 loss_rep=3.6879e+04 total=4.331e+04
it 1/300 lr=0.4833 loss_attr=4.2987e+04 loss_rep=4.5246e+03 total=1.551e+04
it 2/300 lr=0.4667 loss_attr=4.2987e+04 loss_rep=4.5246e+03 total=1.551e+04
it 3/300 lr=0.4500 loss_attr=6.1852e+04 loss_rep=7.5877e+03 total=1.935e+04
it 4/300 lr=0.4333 loss_attr=5.9398e+04 loss_rep=7.1157e+03 total=1.645e+04
it 5/300 lr=0.4167 loss_attr=5.9398e+04 loss_rep=7.1157e+03 total=1.645e+04
it 6/300 lr=0.4000 loss_attr=5.6606e+04 loss_rep=6.8719e+03 total=1.3478e+04
it 7/300 lr=0.3833 loss_attr=5.6606e+04 loss_rep=6.8719e+03 total=1.3478e+04
it 8/300 lr=0.3667 loss_attr=5.3822e+04 loss_rep=6.9836e+03 total=1.0895e+04
it 9/300 lr=0.3500 loss_attr=5.4849e+04 loss_rep=7.8371e+03 total=1.1886e+04
it 10/300 lr=0.3333 loss_attr=5.4849e+04 loss_rep=7.8371e+03 total=1.1886e+04
it 11/300 lr=0.3167 loss_attr=5.1366e+04 loss_rep=7.2349e+03 total=0.8801e+04
it 12/300 lr=0.3000 loss_attr=5.0463e+04 loss_rep=7.3182e+03 total=0.7781e+04
it 13/300 lr=0.2833 loss_attr=4.7822e+04 loss_rep=7.2349e+03 total=0.7781e+04
it 14/300 lr=0.2667 loss_attr=4.7822e+04 loss_rep=7.3365e+03 total=0.4518e+04
it 15/300 lr=0.2500 loss_attr=4.7753e+04 loss_rep=7.7875e+03 total=0.2248e+04
it 16/300 lr=0.2333 loss_attr=4.7753e+04 loss_rep=7.7875e+03 total=0.2248e+04
it 17/300 lr=0.2167 loss_attr=4.2226e+04 loss_rep=7.9646e+03 total=0.10875e+04
it 18/300 lr=0.2000 loss_attr=4.8126e+04 loss_rep=8.1516e+03 total=0.18997e+04
it 19/300 lr=0.1833 loss_attr=4.8126e+04 loss_rep=8.1516e+03 total=0.18997e+04
it 20/300 lr=0.1667 loss_attr=3.7377e+04 loss_rep=8.5566e+03 total=0.145933e+04
it 21/300 lr=0.1500 loss_attr=3.7377e+04 loss_rep=8.5566e+03 total=0.145933e+04
it 22/300 lr=0.1333 loss_attr=3.0884e+04 loss_rep=9.8664e+03 total=0.28712e+04
it 23/300 lr=0.1167 loss_attr=3.1316e+04 loss_rep=9.2139e+03 total=0.45292e+04
it 24/300 lr=0.0800 loss_attr=2.6772e+04 loss_rep=9.0872e+03 total=0.65432e+04
it 25/300 lr=0.0533 loss_attr=2.6772e+04 loss_rep=9.0872e+03 total=0.65432e+04
it 26/300 lr=0.0267 loss_attr=2.2985e+04 loss_rep=9.0418e+03 total=1.3416e+04
it 27/300 lr=0.0000 loss_attr=1.7207e+04 loss_rep=9.0418e+03 total=1.3416e+04
it 28/300 lr=0.8333 loss_attr=1.8405e+04 loss_rep=1.1074e+04 total=0.18079e+04
it 29/300 lr=0.0167 loss_attr=1.7207e+04 loss_rep=1.3899e+04 total=1.0245e+04
it 30/300 lr=0.0000 loss_attr=1.7207e+04 loss_rep=1.4148e+04 total=1.0245e+04
Optimization time (s): 0.4959100100189809
```

```
In [12]: # 8) Plot final embedding colored by labels (if available)
plt.figure(figsize=[8, 6])
if labels is not None:
    plt.scatter(Y_opt[:, 0], Y_opt[:, 1], s=0, cmap='Spectral', alpha=0.8)
else:
    # sc = plt.scatter(sc, Y_opt[:, 0], Y_opt[:, 1], c=labels, s=0, cmap='Spectral', alpha=0.9)
    plt.colorbar(sc, label='label')
    plt.title('Simplified tSNE-like 2D embedding (force-directed)')
    plt.xlabel('Y1')
    plt.ylabel('Y2')
    plt.grid(False)
    plt.tight_layout()
plt.show()
```

(b)

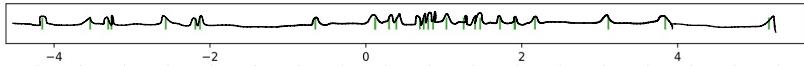
In contrast to last week, we can now clearly distinguish one of the particle types, while the others are still mixed up.



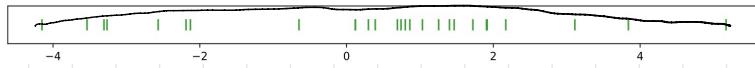
Very good 8/8

② (a) Bandwidth very small: Noisy result, every point has its own small peak

3/3



Bandwidth too large: Flat/smooth curve \rightarrow Some structures can not be seen



For this data set: bandwidth $\approx 0,5$

(b) Depending on the density of points, a smaller bandwidth should be used for regions of higher density and a larger bandwidth in less dense regions.

Use equations to give the proposed method

1/3

③ (a) (y_n, x_n)

$$N = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x-x_0)^2}{2\sigma^2}\right)$$

$$y_n = \beta^T x_n + \epsilon_n$$

$$\hat{\beta} = \arg\max_{\beta} \sum_{n=1}^N \log N(y_n | \beta^T x, \sigma^2)$$

$$\hat{\sigma}^2 = \arg\max_{\sigma^2} \sum_{n=1}^N \log N(y_n | \hat{\beta}^T x, \sigma^2)$$

$$\hat{\beta} = \arg\max_{\beta} \sum_{n=1}^N \log\left(\frac{1}{\sqrt{2\pi}\sigma} \cdot \exp\left(-\frac{(y_n - \beta^T x)^2}{2\sigma^2}\right)\right)$$

$$= \arg\min_{\beta} \sum_{n=1}^N (y_n - \beta^T x_n)^2 = (\vec{y} - X\beta)^T (\vec{y} - X\beta)$$

$$\rightarrow 0 \stackrel{!}{=} \nabla_{\beta} [(\vec{y} - X\beta)^T (\vec{y} - X\beta)]$$

$$= \nabla_{\beta} [-\vec{y}^T X\beta - (\vec{y}^T X\beta) + (\vec{y}^T X\beta)^T X\beta]$$

$$= -X^T \vec{y} - X^T \vec{y} + 2X^T X\beta$$

$$\Rightarrow X^T X\beta = X^T \vec{y}$$

$$\Leftrightarrow \beta = (X^T X)^{-1} X^T \vec{y} \quad \leftarrow \text{Same as result from lecture}$$

$$\hat{\sigma}^2 = \arg\max_{\sigma^2} \sum_{n=1}^N \log\left(\frac{1}{\sqrt{2\pi}\sigma} \cdot \exp\left(-\frac{(y_n - \beta^T x)^2}{2\sigma^2}\right)\right)$$

$$= \arg\max_{\sigma^2} \left[N \log\left(\frac{1}{\sqrt{2\pi}\sigma}\right) - \sum_{n=1}^N \frac{(y_n - \beta^T x)^2}{2\sigma^2} \right]$$

$$= \arg\max_{\sigma^2} \left[-\frac{N}{2} \log(\sigma^2) - \frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - \beta^T x)^2 \right]$$

$$0 \stackrel{!}{=} \frac{\partial}{\partial \sigma^2} \left[\dots \right] = -\frac{N}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{n=1}^N \dots$$

$$N = \frac{1}{\sigma^2} \sum_{n=1}^N (y_n - \beta^T x)^2$$

$$\sigma^2 = \frac{1}{N} \sum_{n=1}^N (y_n - \beta^T x)^2$$

(b) $y_n = \beta^T x_n + \epsilon_n$

$\mathbb{E}(\epsilon_n) = 0, \text{ var}(\epsilon_n) = \sigma_n^2 \leftarrow \text{Now } \sigma_n, \text{ not } \sigma$

$$\hat{\beta} = \arg\max_{\beta} \sum_{n=1}^N \log\left(\frac{1}{\sqrt{2\pi}\sigma_n} \cdot \exp\left(-\frac{(y_n - \beta^T x)^2}{\sigma_n^2}\right)\right)$$

$$= \arg\min_{\beta} \sum_{n=1}^N \frac{(y_n - \beta^T x)^2}{\sigma_n^2}$$

\rightarrow every command is weighted by $\frac{1}{\sigma_n^2}$

$$\boxed{B} \alpha^T \beta = \alpha$$

$$\vec{y}^T X = (X^T \vec{y})^T$$

$$\rightarrow \underset{\beta}{\operatorname{argmin}} (\vec{y} - X\beta)^T W (\vec{y} - X\beta)$$

includes weighting factors

$$0 = \nabla_{\beta} \left[-\vec{y}^T W X \beta - (X\beta)^T W \vec{y} + (X\beta)^T W X \beta \right]$$
$$-X^T W^T \vec{y} - X^T W \vec{y} + 2X^T W X \beta$$

$$\rightarrow (X^T W X) \beta = X^T W \vec{y}$$

$$\beta = (X^T W X)^{-1} X^T W \vec{y}$$

Excellent, 6/6

W is symmetric
 $\begin{pmatrix} \sigma_1^{-2} & & \\ & \ddots & \\ & & \sigma_N^{-2} \end{pmatrix} \Rightarrow W^T = W$